

Project Description (AI_Capstone_Project):

""

DESCRIPTION

Problem Statement

- Amazon is an online shopping website that now caters to millions of people everywhere. Over 34,000 consumer reviews for Amazon brand products like Kindle, Fire TV Stick and more are provided.
- The dataset has attributes like brand, categories, primary categories, reviews.title, reviews.text, and the sentiment. Sentiment is a categorical variable with three levels "Positive", "Negative", and "Neutral". For a given unseen data, the sentiment needs to be predicted.
- You are required to predict Sentiment or Satisfaction of a purchase based on multiple features and review text.

Project Task: Week 1

Class Imbalance Problem:

1. Perform an EDA on the dataset.
 - a) See what a positive, negative, and neutral review looks like
 - b) Check the class count for each class. It's a class imbalance problem.
2. Convert the reviews in Tf-Idf score.
3. Run multinomial Naive Bayes classifier. Everything will be classified as positive because of the class imbalance.

Project Task: Week 2

Tackling Class Imbalance Problem:

1. Oversampling or undersampling can be used to tackle the class imbalance problem.
2. In case of class imbalance criteria, use the following metrics for evaluating model performance: precision, recall, F1-score, AUC-ROC curve. Use F1-Score as the evaluation criteria for this project.
3. Use Tree-based classifiers like Random Forest and XGBoost.

Note: Tree-based classifiers work on two ideologies namely, Bagging or Boosting and have fine-tuning parameter which takes care of the imbalanced class.

Project Task: Week 3

Model Selection:

1. Apply multi-class SVM's and neural nets.
2. Use possible ensemble techniques like: XGboost + oversampled_multinomial_NB.
3. Assign a score to the sentence sentiment (engineer a feature called sentiment score). Use this engineered feature in the model and check for improvements. Draw insights on the same.

Project Task: Week 4

Applying LSTM:

1. Use LSTM for the previous problem (use parameters of LSTM like top-word, embedding-length, Dropout, epochs, number of layers, etc.)

Hint: Another variation of LSTM, GRU (Gated Recurrent Units) can be tried as well.

2. Compare the accuracy of neural nets with traditional ML based algorithms.

3. Find the best setting of LSTM (Neural Net) and GRU that can best classify the reviews as positive, negative, and neutral.

Hint: Use techniques like Grid Search, Cross-Validation and Random Search

Optional Tasks: Week 4

Topic Modeling:

1. Cluster similar reviews.

Note: Some reviews may talk about the device as a gift-option. Other reviews may be about product looks and some may highlight about its battery and performance. Try naming the clusters.

2. Perform Topic Modeling

Hint: Use scikit-learn provided Latent Dirichlette Allocation (LDA) and Non-Negative Matrix Factorization (NMF).

"""

Source Code:

```
import warnings

warnings.filterwarnings('ignore')

import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

import re

import string

import nltk

import seaborn as sns

from sklearn.dummy import DummyClassifier

from sklearn.metrics import precision_score, recall_score, confusion_matrix

from sklearn.metrics import f1_score, roc_auc_score, roc_curve

from sklearn.model_selection import train_test_split, GridSearchCV

from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer

from sklearn.naive_bayes import BernoulliNB, MultinomialNB

from sklearn.linear_model import LogisticRegression

from sklearn.ensemble import RandomForestClassifier

from sklearn import metrics

from sklearn.metrics import roc_auc_score, accuracy_score
```

```
from sklearn.pipeline import Pipeline
```

```
from bs4 import BeautifulSoup
```

```
import re
```

```
import nltk
```

```
from nltk.corpus import stopwords
```

```
from nltk.stem.porter import PorterStemmer
```

```
from nltk.stem import SnowballStemmer, WordNetLemmatizer
```

```
from nltk import sent_tokenize, word_tokenize, pos_tag
```

```
import logging
```

```
from gensim.models import word2vec
```

```
from gensim.models.keyedvectors import KeyedVectors
```

```
from gensim.models import Word2Vec
```

```
from keras.preprocessing import sequence
```

```
from keras.utils import np_utils
```

```
from keras.models import Sequential
```

```
from keras.layers.core import Dense, Dropout, Activation, Lambda
```

```
from keras.layers.embeddings import Embedding
```

```
from keras.layers.recurrent import LSTM, SimpleRNN, GRU
```

```
from keras.preprocessing.text import Tokenizer

from collections import defaultdict

from keras.layers.convolutional import Convolution1D

from keras import backend as K

from keras.layers.embeddings import Embedding

from keras.callbacks import EarlyStopping
```

```
def preprocess(document):

    document = document.lower() # Convert to lowercase

    words = tokenizer.tokenize(document) # Tokenize

    words = [w for w in words if not w in stop_words] # Removing stopwords

    # Lemmatizing

    for pos in [wordnet.NOUN, wordnet.VERB, wordnet.ADJ, wordnet.ADV]:

        words = [wordnet_lemmatizer.lemmatize(x, pos) for x in words]

    return " ".join(words)
```

```
def textPreprocessing(data2):

    #Remove Punctuation Logic
```

```

import string

removePunctuation = [char for char in data2 if char not in string.punctuation]

#Join Chars to form sentences

sentenceWithoutPunctuations = ''.join(removePunctuation)

words = sentenceWithoutPunctuations.split()

#StopwordRemoval

from nltk.corpus import stopwords

removeStopwords = [word for word in words if word.lower() not in
stopwords.words('english')]

return removeStopwords

def cleanText(raw_text, remove_stopwords=False, stemming=False,
split_text=False, \

    ):
'''
Convert a raw review to a cleaned review
'''

text = BeautifulSoup(raw_text, 'lxml').get_text() #remove html

letters_only = re.sub("[^a-zA-Z]", " ", text) # remove non-character

words = letters_only.lower().split() # convert to lower case

```

```
if remove_stopwords: # remove stopword  
    stops = set(stopwords.words("english"))  
    words = [w for w in words if not w in stops]
```

```
if stemming==True: # stemming  
    stemmer = SnowballStemmer('english')  
    words = [stemmer.stem(w) for w in words]
```

```
if split_text==True: # split text  
    return (words)
```

```
return( " ".join(words))
```

```
def modelEvaluation(predictions):  
    '''  
    Print model evaluation to predicted result  
    '''  
  
    print ("\nAccuracy on validation set: {:.4f}".format(accuracy_score(y_test,  
predictions)))  
  
    #print("\nAUC score : {:.4f}".format(roc_auc_score(y_test, predictions)))
```

```
print("\nClassification report : \n", metrics.classification_report(y_test,
predictions))

print("\nConfusion Matrix : \n", metrics.confusion_matrix(y_test, predictions))
```

```
def parseSent(review, tokenizer, remove_stopwords=False):
    """
    Parse text into sentences
    """
    raw_sentences = tokenizer.tokenize(review.strip())
    sentences = []
    for raw_sentence in raw_sentences:
        if len(raw_sentence) > 0:
            sentences.append(cleanText(raw_sentence, remove_stopwords,
split_text=True))
    return sentences
```

```
def makeFeatureVec(review, model, num_features):
    """
```


Transform a review to a feature vector by averaging feature vectors of words appeared in that review and in the vocabulary list created

```
'''
```

```
featureVec = np.zeros((num_features,),dtype="float32")
```

```
nwords = 0.
```

```
index2word_set = set(model.wv.index_to_key) #index2word is the vocabulary list of the Word2Vec model
```

```
isZeroVec = True
```

```
for word in review:
```

```
    if word in index2word_set:
```

```
        nwords = nwords + 1.
```

```
        featureVec = np.add(featureVec, model.wv[word])
```

```
        isZeroVec = False
```

```
if isZeroVec == False:
```

```
    featureVec = np.divide(featureVec, nwords)
```

```
return featureVec
```

```
def getAvgFeatureVecs(reviews, model, num_features):
```

```
'''
```

```
Transform all reviews to feature vectors using makeFeatureVec()
```

```
'''

counter = 0

reviewFeatureVecs = np.zeros((len(reviews),num_features),dtype="float32")

for review in reviews:

    reviewFeatureVecs[counter] = makeFeatureVec(review,
model,num_features)

    counter = counter + 1

return reviewFeatureVecs
```

#Reading Data sheets

```
file_path=input("enter path for the loan data file to load:")

df_path=file_path.replace("\\","/")

data = pd.read_csv(df_path)
```

```
file_path=input("enter path for the loan data file to load:")

df_path=file_path.replace("\\","/")
```

```
test = pd.read_csv(df_path)
```

```
file_path=input("enter path for the loan data file to load:")
```

```
df_path=file_path.replace("\\",'/')
```

```
test_prediction = pd.read_csv(df_path)
```

```
print(data.head())
```

```
#WEEK1
```

```
#See what a positive, negative, and neutral review looks like
```

```
Positive = data[data['sentiment']== "Positive"].iloc[:,[5,6,7]]
```

```
Neutral = data[data['sentiment']== "Neutral"].iloc[:,[5,6,7]]
```

```
Negative = data[data['sentiment']== "Negative"].iloc[:,[5,6,7]]
```

```
print("See what a positive, negative, and neutral review looks like:")
```

```
positive=data[['sentiment']]=="Positive"
```

```
sns.distplot(positive['sentiment'])
```

```
plt.title("Positive Reviews")
```

```
plt.show()
```

```
negative=data[['sentiment']] == "Negative"
```

```
sns.distplot(negative['sentiment'])
```

```
plt.title("Negative Reviews")
```

```
plt.show()
```

```
neutral=data[['sentiment']] == "Neutral"
```

```
sns.distplot(neutral['sentiment'])
```

```
plt.title("Neutral Reviews")
```

```
plt.show()
```

```
print("_____")  
_____")
```

```
#Check the class count for each class. It's a class imbalance problem
```

```
print("Check the class count for each class")
```

```
print("-----\n")
```

```
print(Positive['sentiment'].value_counts())
```

```
print(Neutral['sentiment'].value_counts())
```

```
print(Negative['sentiment'].value_counts())
```

```
print("_____")
```

```
# Keeping only those Features that we need for further exploring.
```

```
data1 = data [['sentiment',"reviews.text"]]
```

```
# Resetting the Index.
```

```
data1.index = pd.Series(list(range(data1.shape[0])))
```

```
from nltk.tokenize import RegexpTokenizer
```

```
from nltk.corpus import stopwords
```

```
import nltk
```

```
from nltk.corpus import wordnet
```

```
from nltk.stem import WordNetLemmatizer
```

```
nltk.download('wordnet')
```

```
#Download Stopwords
```

```
nltk.download('stopwords')
```

```
wordnet_lemmatizer = WordNetLemmatizer()
```

```
tokenizer = RegexpTokenizer(r'[a-z]+')
```

```
stop_words = set(stopwords.words('english'))
```

```
data1['Processed_Review'] = data1['reviews.text'].apply(preprocess)
```

```
data2 = data1[["sentiment", "Processed_Review"]]
```

```
print(data2.groupby('sentiment').describe())
```

```
#Text preprocessing
```

```
data2['Processed_Review'].head(2).apply(textPreprocessing)
```

```
from sklearn.feature_extraction.text import CountVectorizer
```

```
bow =
```

```
CountVectorizer(analyzer=textPreprocessing).fit(data2['Processed_Review'])
```

```
reviews_bow = bow.transform(data2['Processed_Review'])
```

```
print("_____")
```

```
#Convert the reviews in Tf-Idf score.
```

```
from sklearn.feature_extraction.text import TfidfTransformer
```

```
tfidfData = TfidfTransformer().fit(reviews_bow)
tfidfDataFinal = tfidfData.transform(reviews_bow)
print("Convert the reviews in Tf-Idf score:\n")
print(tfidfDataFinal)
```

```
print("_____")
_____")
```

#Run multinomial Naive Bayes classifier. Everything will be classified as positive because of the class imbalance.

```
print("Run multinomial Naive Bayes classifier. Everything will be classified as positive because of the class imbalance.")
```

```
from sklearn.naive_bayes import MultinomialNB
```

```
model = MultinomialNB().fit(tfidfDataFinal,data2['sentiment'])
```

```
inputData = "very bad dont like it at all it sucks"
```

```
l1 = textPreprocessing(inputData)
```

```
l2 = bow.transform(l1)
```

```
l3 = tfidfData.transform(l2)
```

```
prediction = model.predict(l3[0])
```

```
print(prediction)
```

```
print("_____")
_____")
```

#Creating independent and Dependent Features

```
columns = data2.columns.tolist()
```

Filtering the columns to remove data we do not want

```
columns = [c for c in columns if c not in ["sentiment"]]
```

Store the variable we are predicting

```
target = "sentiment"
```

Defining a random state

```
state = np.random.RandomState(42)
```

```
X = data2[columns]
```

```
Y = data2[target]
```

#WEEK2

#Oversampling or undersampling can be used to tackle the class imbalance problem

```
print("Oversampling or undersampling can be used to tackle the class imbalance problem")
```

RandomOverSampler to handle imbalanced data

```
from imblearn.over_sampling import RandomOverSampler
```

```
ros = RandomOverSampler(random_state=0)
```

```
X_res,Y_res=ros.fit_resample(X,Y)
```



```
from collections import Counter

print(sorted(Counter(Y_res).items()))
```

```
#Checking out both old & new data

print('Original dataset shape {}'.format(Counter(Y)))

print('Resampled dataset shape {}'.format(Counter(Y_res)))
```

```
#Creating X output to dataframe

X1=pd.DataFrame(X_res,columns=['Processed_Review'])
```

```
#Creating Y output to dataframe for merging

Y1=pd.DataFrame(Y_res,columns=['sentiment'])
```

```
#Merging the X & Y output to Final data
```

```
Final_data=pd.concat([X1,Y1],axis=1)
```

```
print(Final_data.head())
```

```
df = Final_data.sample(frac=0.1, random_state=0)
```

```
# Dropping missing values
```

```
df.dropna(inplace=True)
```

```
df.head()
```

```
print("_____")  
_____")
```

```
# Splitting data into training set and validation
```

```
X_train, X_test, y_train, y_test = train_test_split(df['Processed_Review'],  
df['sentiment'], \
```

```
test_size=0.1, random_state=0)
```

```
y_tra=y_train
```

```
# Preprocess text data in training set and validation set
```

```
X_train_cleaned = []
```

```
X_test_cleaned = []
```

```
for d in X_train:
```

```
    X_train_cleaned.append(cleanText(d))
```

```
for d in X_test:
```

```
    X_test_cleaned.append(cleanText(d))
```

```
# Fit and transform the training data to a document-term matrix using  
CountVectorizer
```

```
countVect = CountVectorizer()
```

```
X_train_countVect = countVect.fit_transform(X_train_cleaned)
```

```
# Train MultinomialNB classifier
```

```
mnb = MultinomialNB()
```

```
mnb.fit(X_train_countVect, y_train)
```

```
print("MultinomialNB classifier")
```

```
predictions = mnb.predict(countVect.transform(X_test_cleaned))
```

```
modelEvaluation(predictions)
```

```
print("_____")
```

```
print("XGBoost Classifier")
```

```
print("-----\n")
```

```
from xgboost import XGBClassifier
```

```
# Fitting and transforming the training data to a document-term matrix using  
TfidfVectorizer
```

```
tfidf = TfidfVectorizer(min_df=5) #minimum document frequency of 5
```

```
X_train_tfidf = tfidf.fit_transform(X_train)

print("Number of features : %d \n" %len(tfidf.get_feature_names())) #1722

print("Show some feature names : \n", tfidf.get_feature_names()[::1000])
```

```
# XGBoost Classifier
```

```
xgb = XGBClassifier()

xgb.fit(X_train_tfidf, y_train)
```

```
# Evaluating on the validation set
```

```
predictions = xgb.predict(tfidf.transform(X_test_cleaned))

modelEvaluation(predictions)
```

```
print("_____")
print("_____")
```

```
sentences = []
```

```
for review in X_train_cleaned:
```

```
    sentences += parseSent(review, tokenizer)
```

```
from gensim.models import Word2Vec
```

```
w2v = Word2Vec()
```

```
# Fitting parsed sentences to Word2Vec model
```

```
num_features = 300 #embedding dimension
```

```
min_word_count = 10
```

```
num_workers = 4
```

```
context = 10
```

```
downsampling = 1e-3
```

```
w2v = Word2Vec(sentences, workers=num_workers, vector_size=num_features,  
min_count = min_word_count,\
```

```
    window = context, sample = downsampling)
```

```
w2v.init_sims(replace=True)
```

```
w2v.save("w2v_300features_10minwordcounts_10context") #save trained  
word2vec model
```

```
X_train_cleaned1 = []
```

```
for review in X_train:
```

```
    X_train_cleaned1.append(cleanText(review, remove_stopwords=True,  
split_text=True))
```

```
trainVector = getAvgFeatureVecs(X_train_cleaned1, w2v, num_features)
```

```
# Getting feature vectors for validation set
```

```
X_test_cleaned1 = []
```

```
for review in X_test:
```

```
    X_test_cleaned1.append(cleanText(review, remove_stopwords=True,  
split_text=True))
```

```
testVector = getAvgFeatureVecs(X_test_cleaned1, w2v, num_features)
```

```
# Getting feature vectors for training set
```

```
trainVector = getAvgFeatureVecs(X_train, w2v, num_features)
```

```
# Getting feature vectors for validation set
```

```
testVector = getAvgFeatureVecs(X_test, w2v, num_features)
```

```
# Random Forest Classifier
```

```
print("Random Forest Classifier")
```

```
print("-----\n")
```

```
rf = RandomForestClassifier(n_estimators=100)
```

```
rf.fit(trainVector, y_train)
```

```
predictions = rf.predict(testVector)
```

```
modelEvaluation(predictions)
```

```
df = Final_data.sample(frac=0.1, random_state=0)
```

```
# Drop missing values
```

```
df.dropna(inplace=True)
```

```
print("_____")  
_____")
```

```
#Apply multi-class SVM's and neural nets.
```

```
print("Apply multi-class SVM's and neural nets.")
```

```
print("-----")
```

```
# Fitting and transforming the training data to a document-term matrix using  
TfidfVectorizer
```

```
tfidf = TfidfVectorizer(min_df=5) #minimum document frequency of 5
```

```
X_train_tfidf = tfidf.fit_transform(X_train)
```

```
# Logistic Regression
```

```
print("Logistic Regression")
```

```
lr = LogisticRegression()
```

```
lr.fit(X_train_tfidf, y_train)
```

```
# Have a look at the top 10 features with the smallest and largest coefficients
```

```
feature_names = np.array(tfidf.get_feature_names())
```

```
sorted_coef_index = lr.coef_[0].argsort()
```

```
# Evaluating on the validation set
```

```
predictions = lr.predict(tfidf.transform(X_test_cleaned))
```

```
modelEvaluation(predictions)
```

```
print("_____")
```

```
# Fitting and transforming the training data to a document-term matrix using  
TfidfVectorizer
```

```
tfidf = TfidfVectorizer(min_df=5) #minimum document frequency of 5
```

```
X_train_tfidf = tfidf.fit_transform(X_train)
```

```
#Apply multi-class SVM's and neural nets.
```



```

# SVM

from sklearn.linear_model import SGDClassifier

clf = SGDClassifier(loss="hinge", penalty="l2")

clf.fit(X_train_tfidf, y_train)

print("SGDClassifier")

print("-----\n")

# Have a look at the top 10 features with the smallest and largest coefficients

feature_names = np.array(tfidf.get_feature_names())

sorted_coef_index = clf.coef_[0].argsort()


# Evaluating on the validation set

predictions = clf.predict(tfidf.transform(X_test_cleaned))

modelEvaluation(predictions)


print("_____")
print("_____")

from xgboost import XGBClassifier

# Fitting and transforming the training data to a document-term matrix using
TfidfVectorizer

tfidf = TfidfVectorizer(min_df=5) #minimum document frequency of 5

X_train_tfidf = tfidf.fit_transform(X_train)

```

```
# XGBoost Classifier
```

```
print("XGBoost Classifier")
```

```
print("-----")
```

```
xgb = XGBClassifier()
```

```
xgb.fit(X_train_tfidf, y_train)
```

```
# Look at the top 10 features with smallest and the largest coefficients
```

```
feature_names = np.array(tfidf.get_feature_names())
```

```
# sorted_coef_index = xgb.coef_[0].argsort()
```

```
# Evaluating on the validation set
```

```
predictions = xgb.predict(tfidf.transform(X_test_cleaned))
```

```
modelEvaluation(predictions)
```

```
print("_____")
```

```
#Assign a score to the sentence sentiment
```

```
print("Assign a score to the sentence sentiment :\n")
```

```
print("-----")
```

```
# Convert the sentiments
```

```
df.sentiment.replace(('Positive','Negative','Neutral'),(1,0,2),inplace=True)
```

```
print(df.head())
```

```
print("_____")
```

```
# Splitting data into training set and validation
```

```
X_train, X_test, y_train, y_test = train_test_split(df['Processed_Review'],  
df['sentiment'], \
```

```
test_size=0.1, random_state=1)
```

```
top_words = 20000
```

```
maxlen = 100
```

```
batch_size = 32
```

```
nb_classes = 3
```

```
nb_epoch = 3
```

```
# Vectorize X_train and X_test to 2D tensor
```

```
tokenizer = Tokenizer(nb_words=top_words) #Considering only top 20000 words  
in the corpus
```

```
tokenizer.fit_on_texts(X_train)
```

```
# tokenizer.word_index #access word-to-index dictionary of trained tokenizer
```

```
sequences_train = tokenizer.texts_to_sequences(X_train)
```

```
sequences_test = tokenizer.texts_to_sequences(X_test)
```

```
X_train_seq = sequence.pad_sequences(sequences_train, maxlen=maxlen)
```

```
X_test_seq = sequence.pad_sequences(sequences_test, maxlen=maxlen)
```

```
# One-Hot Encoding of y_train and y_test
```

```
y_train_seq = np_utils.to_categorical(y_train, nb_classes)
```

```
y_test_seq = np_utils.to_categorical(y_test, nb_classes)
```

```
# Constructing a Simple LSTM
```

```
print("Constructing a Simple LSTM")
```

```
print("-----\n")

model1 = Sequential()

model1.add(Embedding(top_words, 128))

model1.add(Dropout(0.2))

model1.add(LSTM(128))

model1.add(Dropout(0.2))

model1.add(Dropout(0.2))

model1.add(Dense(nb_classes))

model1.add(Activation('softmax'))

model1.summary()


# Compiling LSTM

model1.compile(loss='binary_crossentropy',
               optimizer='adam',
               metrics=['accuracy'])


model1.fit(X_train_seq, y_train_seq, batch_size=batch_size, epochs=nb_epoch,
          verbose=1)


# Model Evaluation

score = model1.evaluate(X_test_seq, y_test_seq, batch_size=batch_size)
```

```

print('Test loss : {:.4f}'.format(score[0]))

print('Test accuracy : {:.4f}'.format(score[1]))


# Getting weight matrix of the embedding layer

model1.layers[0].get_weights()[0] # weight matrix of the embedding layer, word-
by-dim matrix

print("Size of weight matrix in the embedding layer : ", \
      model1.layers[0].get_weights()[0].shape) #(20000, 128)


# Getting weight matrix of the hidden layer

print("Size of weight matrix in the hidden layer : ", \
      model1.layers[2].get_weights()[0].shape) #(128, 512) weight dim of LSTM - w


# Getting weight matrix of the output layer

print("Size of weight matrix in the output layer : ", \
      model1.layers[5].get_weights()[0].shape) #(128, 2) weight dim of dense layer


# Loading pretrained Word2Vec model

w2v = Word2Vec.load("w2v_300features_10minwordcounts_10context")

```

```
# Getting Word2Vec embedding matrix
```

```
embedding_matrix = w2v.wv.vectors # embedding matrix, type = numpy.ndarray
```

```
print("Shape of embedding matrix : ", embedding_matrix.shape) #(4016, 300) =  
(vocabulary size, embedding dimension)
```

```
# w2v.wv.syn0[0] #feature vector of the first word in the vocabulary list
```

```
top_words = embedding_matrix.shape[0] #4016
```

```
maxlen = 100
```

```
batch_size = 32
```

```
nb_classes = 3
```

```
nb_epoch = 3
```

```
# Vectorizing X_train and X_test to 2D tensor
```

```
tokenizer = Tokenizer(nb_words=top_words) #Considering only top 20000 words  
in the corpus
```

```
tokenizer.fit_on_texts(X_train)
```

```
# tokenizer.word_index #access word-to-index dictionary of trained tokenizer
```

```
sequences_train = tokenizer.texts_to_sequences(X_train)
```

```
sequences_test = tokenizer.texts_to_sequences(X_test)
```

```
X_train_seq = sequence.pad_sequences(sequences_train, maxlen=maxlen)
```

```
X_test_seq = sequence.pad_sequences(sequences_test, maxlen=maxlen)
```

```
# One-Hot Encoding of y_train and y_test
```

```
y_train_seq = np_utils.to_categorical(y_train, nb_classes)
```

```
y_test_seq = np_utils.to_categorical(y_test, nb_classes)
```

```
print('X_train shape:', X_train_seq.shape) #(27799, 100)
```

```
print('X_test shape:', X_test_seq.shape) #(3089, 100)
```

```
print('y_train shape:', y_train_seq.shape) #(27799, 2)
```

```
print('y_test shape:', y_test_seq.shape) #(3089, 2)
```

```
# Constructing Word2Vec embedding layer
```

```
embedding_layer = Embedding(embedding_matrix.shape[0], #4016
```

```
    embedding_matrix.shape[1], #300
```

```
    weights=[embedding_matrix])
```



```
print("_____")
```

```
# Constructing LSTM with Word2Vec embedding
```

```
print("Constructing LSTM with Word2Vec embedding")
```

```
print("-----\n")
```

```
model2 = Sequential()
```

```
model2.add(embedding_layer)
```

```
model2.add(LSTM(128))
```

```
model2.add(Dropout(0.2))
```

```
model2.add(Dropout(0.2))
```

```
model2.add(Dense(nb_classes))
```

```
model2.add(Activation('softmax'))
```

```
model2.summary()
```

```
# Compiling model
```

```
model2.compile(loss='binary_crossentropy',
```

```
               optimizer='adam',
```

```
               metrics=['accuracy'])
```

```
model2.fit(X_train_seq, y_train_seq, batch_size=batch_size, epochs=nb_epoch,  
verbose=1)
```

```
# Model evaluation
```

```
score = model2.evaluate(X_test_seq, y_test_seq, batch_size=batch_size)
```

```
print('Test loss : {:.4f}'.format(score[0]))
```

```
print('Test accuracy : {:.4f}'.format(score[1]))
```

```
# Getting weight matrix of the embedding layer
```

```
print("Size of weight matrix in the embedding layer : ", \
```

```
      model2.layers[0].get_weights()[0].shape) #(20000, 128)
```

```
# Getting weight matrix of the hidden layer
```

```
print("Size of weight matrix in the hidden layer : ", \
```

```
      model2.layers[1].get_weights()[0].shape) #(128, 512) weight dim of LSTM - w
```

```
# Getting weight matrix of the output layer
```

```
print("Size of weight matrix in the output layer : ", \
```

```
      model2.layers[4].get_weights()[0].shape) #(128, 2) weight dim of dense layer
```

```
print("_____")
```

```
print("Find the best setting of LSTM (Neural Net) and GRU that can best classify  
the reviews as positive, negative, and neutral.\nHint: Use techniques like Grid  
Search, Cross-Validation and Random Search")
```

```
print("-----  
-----")
```

```
# Building a pipeline
```

```
estimators = [("tfidf", TfidfVectorizer()), ("lr", LogisticRegression())]
```

```
model = Pipeline(estimators)
```

```
# Grid search
```

```
params = {"lr__C": [0.1, 1, 10], #regularization param of logistic regression
```

```
         "tfidf__min_df": [1, 3], #min count of words
```

```
         "tfidf__max_features": [1000, None], #max features
```

```
         "tfidf__ngram_range": [(1,1), (1,2)], #1-grams or 2-grams
```

```
         "tfidf__stop_words": [None, "english"]} #use stopwords or don't
```

```
grid = GridSearchCV(estimator=model, param_grid=params, scoring="accuracy",  
n_jobs=-1)
```

```
grid.fit(X_train, y_train)
```

```
print("\nGrid search:\n")
```

```
print("The best parameter set is : \n", grid.best_params_)
```

```
predictions = grid.predict(X_test)
```

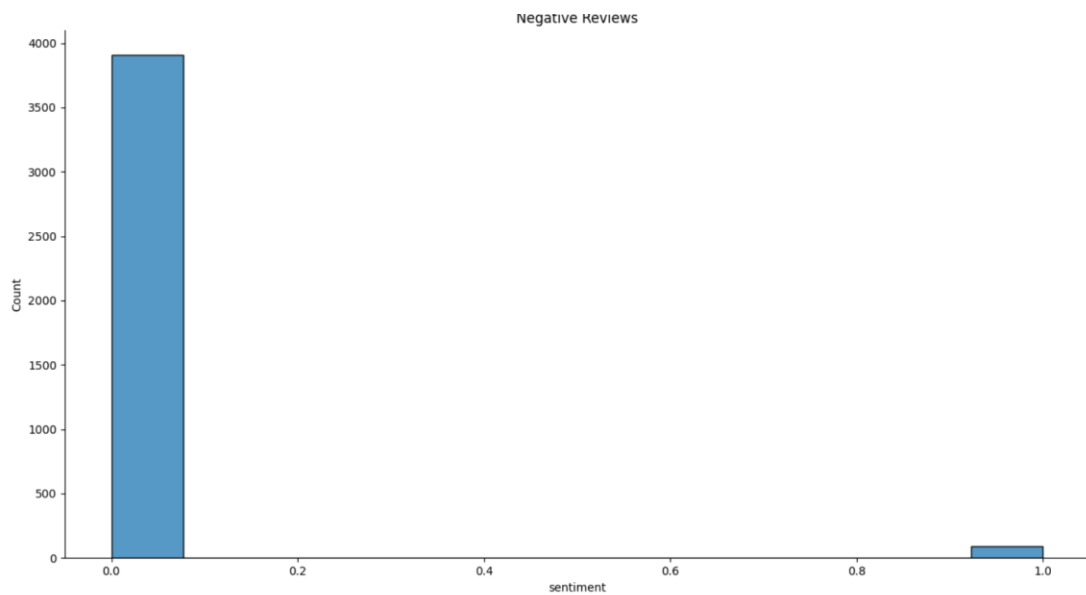
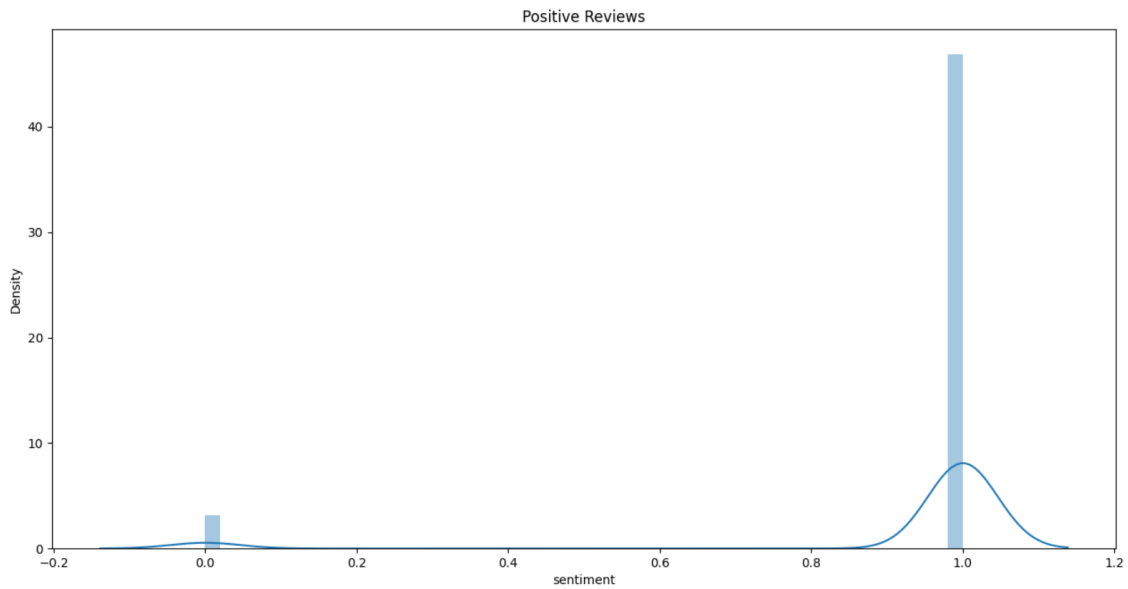
```
modelEvaluation(predictions)
```

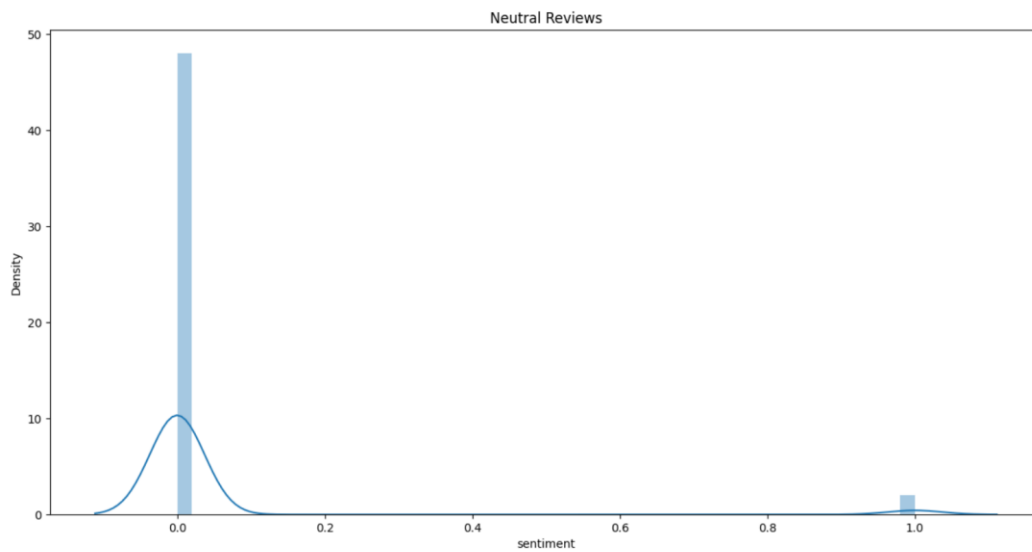
Screenshot of the output:

Week 1:

Task1: Perform an EDA on the dataset.

a) See what a positive, negative, and neutral review looks like





b) Check the class count for each class. It's a class imbalance problem.:

```
Check the class count for each class
```

```
-----
```

```
Positive    3749
Name: sentiment, dtype: int64
Neutral     158
Name: sentiment, dtype: int64
Negative     93
Name: sentiment, dtype: int64
```

Task 2: Convert the reviews in Tf-Idf score:

```
Convert the reviews in Tf-Idf score:
```

```
(0, 3292)    0.12348731897288433
(0, 2955)    0.13262765127937107
(0, 2894)    0.22889129312178152
(0, 2861)    0.22889129312178152
(0, 2782)    0.22889129312178152
(0, 2565)    0.13750705381771683
(0, 2480)    0.18120952660343198
(0, 2361)    0.21044546789025603
(0, 2350)    0.08672269357481559
(0, 2315)    0.17527129722802176
(0, 2292)    0.08954066746770499
(0, 2259)    0.18120952660343198
(0, 2210)    0.09077190449626162
(0, 2139)    0.14545045985637628
(0, 2055)    0.08206248864024848
(0, 1971)    0.09636033823619856
(0, 1922)    0.13279862678760768
(0, 1523)    0.22889129312178152
(0, 1348)    0.12108339495186818
(0, 1319)    0.13174106906037267
(0, 1236)    0.22889129312178152
(0, 1128)    0.11971297627771318
(0, 1062)    0.17527129722802176
(0, 1041)    0.11719661158900058
(0, 678)     0.21044546789025603
:           :
```

Task 3: Run multinomial Naive Bayes classifier. Everything will be classified as positive because of the class imbalance.:

```
Run multinomial Naive Bayes classifier. Everything will be classified as positive because of the class imbalance.
['Positive']
```

WEEK 2:

Task 1: Oversampling or undersampling can be used to tackle the class imbalance problem.:

```
Oversampling or undersampling can be used to tackle the class imbalance problem
[('Negative', 3749), ('Neutral', 3749), ('Positive', 3749)]
Original dataset shape Counter({'Positive': 3749, 'Neutral': 158, 'Negative': 93})
Resampled dataset shape Counter({'Positive': 3749, 'Neutral': 3749, 'Negative': 3749})
      Processed_Review sentiment
0  purchase black fridaypros great price even sal... Positive
1  purchase two amazon echo plus two dot plus fou... Positive
2  average alexa option show thing screen still l... Neutral
3              good product exactly want good price Positive
4  rd one purchase buy one niece case compare one... Positive
```

Task 2: In case of class imbalance criteria, use the following metrics for evaluating model performance: precision, recall, F1-score, AUC-ROC curve. Use F1-Score as the evaluation criteria for this project.:

MultinomialNB classifier

Accuracy on validation set: 0.8938

Classification report :

	precision	recall	f1-score	support
Negative	0.93	0.95	0.94	39
Neutral	0.85	0.90	0.88	39
Positive	0.91	0.83	0.87	35
accuracy			0.89	113
macro avg	0.89	0.89	0.89	113
weighted avg	0.89	0.89	0.89	113

Confusion Matrix :

```
[[37  0  2]
 [ 3 35  1]
 [ 0  6 29]]
```

Task 3: Use Tree-based classifiers like Random Forest and XGBoost.:

```
XGBoost Classifier
-----

Number of features : 691

Show some feature names :
['able']
[08:43:58] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the
from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

Accuracy on validation set: 0.9292

Classification report :
      precision    recall  f1-score   support

   Negative      0.95      0.97      0.96        39
    Neutral      0.90      0.95      0.92        39
    Positive      0.94      0.86      0.90        35

 accuracy          0.93          0.93          0.93          113
 macro avg          0.93          0.93          0.93          113
weighted avg          0.93          0.93          0.93          113

Confusion Matrix :
[[38  0  1]
 [ 1 37  1]
 [ 1  4 30]]
```

```
Random Forest Classifier
-----

Accuracy on validation set: 0.4602

Classification report :
      precision    recall  f1-score   support

   Negative      0.49      0.69      0.57        39
    Neutral      0.44      0.36      0.39        39
    Positive      0.42      0.31      0.36        35

 accuracy          0.46          0.46          0.46          113
 macro avg          0.45          0.46          0.44          113
weighted avg          0.45          0.46          0.45          113

Confusion Matrix :
[[27  6  6]
 [16 14  9]
 [12 12 11]]
```

WEEK 3:

Task 1: Apply multi-class SVM's and neural nets.:

Apply multi-class SVM's and neural nets.

Logistic Regression

Accuracy on validation set: 0.9292

Classification report :

	precision	recall	f1-score	support
Negative	0.93	1.00	0.96	39
Neutral	0.88	0.92	0.90	39
Positive	1.00	0.86	0.92	35
accuracy			0.93	113
macro avg	0.94	0.93	0.93	113
weighted avg	0.93	0.93	0.93	113

Confusion Matrix :

```
[[39  0  0]
 [ 3 36  0]
 [ 0  5 30]]
```

SGDClassifier

Accuracy on validation set: 0.9204

Classification report :

	precision	recall	f1-score	support
Negative	0.93	1.00	0.96	39
Neutral	0.86	0.95	0.90	39
Positive	1.00	0.80	0.89	35
accuracy			0.92	113
macro avg	0.93	0.92	0.92	113
weighted avg	0.93	0.92	0.92	113

Confusion Matrix :

```
[[39  0  0]
 [ 2 37  0]
 [ 1  6 28]]
```


Task 2: Use possible ensemble techniques like: XGboost + oversampled_multinomial_NB.:

```
XGBoost Classifier
-----
[08:44:03] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src
from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore th

Accuracy on validation set: 0.9292

Classification report :
      precision    recall  f1-score   support

   Negative       0.95     0.97     0.96         39
   Neutral       0.90     0.95     0.92         39
   Positive       0.94     0.86     0.90         35

 accuracy          0.93
 macro avg         0.93
weighted avg         0.93

Confusion Matrix :
[[38  0  1]
 [ 1 37  1]
 [ 1  4 30]]
```

Task 3: Assign a score to the sentence sentiment:

```
Assign a score to the sentence sentiment :
-----
      Processed_Review  sentiment
8805  buy think would great read book play game howe...      2
9736      good tablet kid lot appts download game      2
125      item work expect great product      1
10143  great beginner like child limit use many apps ...      2
10937  buy kindle past time one come defective port b...      2
```

WEEK 4:

Task 1: Use LSTM for the previous problem (use parameters of LSTM like top-word, embedding-length, Dropout, epochs, number of layers, etc.):

```
Constructing a Simple LSTM
-----

2022-03-02 08:44:04.153837: W tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load dynamic
2022-03-02 08:44:04.153942: W tensorflow/stream_executor/cuda/cuda_driver.cc:269] failed call to cuInit: UNKNOWN ER
2022-03-02 08:44:04.157280: I tensorflow/stream_executor/cuda/cuda_diagnostics.cc:169] retrieving CUDA diagnostic in
2022-03-02 08:44:04.157576: I tensorflow/stream_executor/cuda/cuda_diagnostics.cc:176] hostname: LAPTOP-TU8FR7UU
2022-03-02 08:44:04.158152: I tensorflow/core/platform/cpu_feature_guard.cc:142] This TensorFlow binary is optimized
ormance-critical operations:  AVX AVX2
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.
Model: "sequential"

Layer (type)                Output Shape                Param #
=====
embedding (Embedding)       (None, None, 128)          2560000
dropout (Dropout)           (None, None, 128)          0
lstm (LSTM)                 (None, 128)                131584
dropout_1 (Dropout)         (None, 128)                0
dropout_2 (Dropout)         (None, 128)                0
dense (Dense)               (None, 3)                  387
activation (Activation)     (None, 3)                  0
=====
Total params: 2,691,971
Trainable params: 2,691,971
Non-trainable params: 0

2022-03-02 08:44:04.648983: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of the MLIR Optimi
Epoch 1/3
32/32 [=====] - 7s 174ms/step - loss: 0.6507 - accuracy: 0.4140
Epoch 2/3
32/32 [=====] - 5s 171ms/step - loss: 0.5617 - accuracy: 0.6709
Epoch 3/3
32/32 [=====] - 6s 174ms/step - loss: 0.2963 - accuracy: 0.8715
4/4 [=====] - 0s 24ms/step - loss: 0.2711 - accuracy: 0.8407
Test loss : 0.2711
Test accuracy : 0.8407
Size of weight matrix in the embedding layer : (20000, 128)
Size of weight matrix in the hidden layer : (128, 512)
Size of weight matrix in the output layer : (128, 3)
Shape of embedding matrix : (416, 300)
X_train shape: (1012, 100)
X_test shape: (113, 100)
```

Task 3: Find the best setting of LSTM (Neural Net) and GRU that can best classify the reviews as positive, negative, and neutral.

Hint: Use techniques like Grid Search, Cross-Validation and Random Search:

```
Find the best setting of LSTM (Neural Net) and GRU that can best classify the reviews as positive, negative, and neutral.  
Hint: Use techniques like Grid Search, Cross-Validation and Random Search
```

```
-----  
Grid search:
```

```
The best parameter set is :
```

```
{'lr_C': 10, 'tfidf__max_features': None, 'tfidf__min_df': 1, 'tfidf__ngram_range': (1, 2), 'tfidf__stop_words': None}
```

```
Accuracy on validation set: 0.9381
```

```
Classification report :
```

	precision	recall	f1-score	support
0	0.97	1.00	0.99	33
1	0.94	0.89	0.91	35
2	0.91	0.93	0.92	45
accuracy			0.94	113
macro avg	0.94	0.94	0.94	113
weighted avg	0.94	0.94	0.94	113

```
Confusion Matrix :
```

```
[[33  0  0]  
[ 0 31  4]  
[ 1  2 42]]
```